

Movie Review Sentiment Analysis with Advanced NLP Techniques

Author: Rimsha Mahmood , CMS: 455080

May 2025

Abstract

This project conducts sentiment analysis on 50,000 IMDb movie reviews, classifying them as positive or negative using Logistic Regression, Multinomial Naive Bayes, and Long Short-Term Memory (LSTM) models. Employing TF-IDF vectorization and LSTM embeddings, the study achieves a peak accuracy of 89.14% with Logistic Regression. Visualizations, including word clouds and ROC curves, elucidate sentiment patterns, while aspect-based analysis explores sentiments toward specific movie aspects (e.g., acting, plot). The report discusses limitations, ethical considerations, and future directions, such as transformer models and explainability tools.

Contents

1	Introduction	2
2	Dataset Description	2
3	Data Exploration	3
4	Methodology	4
4.1	Data Preprocessing	4
4.2	Feature Engineering	5
4.3	Data Splitting	6
4.4	Aspect-Based Sentiment Analysis	6
5	Sentiment Analysis Models	7
5.1	Logistic Regression	7
5.2	Multinomial Naive Bayes	7
5.3	Long Short-Term Memory (LSTM)	8
6	Feature Importance Analysis	9
7	Aspect-Based Sentiment Results	9
8	Results and Discussion	10
8.1	Quantitative Results	10
8.2	Visualization Analysis	10
8.3	Overfitting Assessment	12
9	Limitations	12
10	Ethical Considerations	12
11	Conclusions and Future Work	12
A	Supplementary Code	13

1 Introduction

Sentiment analysis is pivotal for understanding opinions in text data. This project analyzes a 50,000-review IMDb dataset (<https://www.kaggle.com/datasets/lakshmi25npathi/imdb-dataset-of-50k-movie-reviews>), comparing Logistic Regression, Multinomial Naive Bayes, and LSTM models for classifying reviews as positive or negative. It also explores aspect-based sentiment analysis and employs visualizations (e.g., confusion matrices, word clouds) for interpretability.

Objectives:

- Achieve high accuracy in binary sentiment classification.
- Compare traditional and deep learning NLP approaches.
- Analyze sentiments for specific movie aspects.
- Provide clear visualizations of results.

2 Dataset Description

The IMDb dataset comprises 50,000 movie reviews, initially balanced with 25,000 positive and 25,000 negative reviews. After removing 418 duplicates, it contains 49,582 reviews (24,791 positive, 24,791 negative). The dataset has two columns, described in Table 1.

Table 1: Dataset Columns Description

Column	Description
review	The text content of the movie review, ranging from 100 to 2,470 words.
sentiment	The sentiment label, either positive or negative.

Characteristics:

- **Balance:** 25000 positive and 25000 negative reviews.
- **Length:** Mean 231 words (std 171), with negative reviews slightly longer.
- **Quality:** No missing values; duplicates removed.




Figure 1: Sentiment Distribution of IMDb Reviews

3 Data Exploration

Exploratory analysis revealed:

- **Sentiment:** Balanced distribution post-deduplication (25000 positive, 25000 negative).
- **Length:** Negative reviews slightly longer (mean 231 words, std 171).
- **Words:** Post-preprocessing, positive reviews highlight “great”, “movie”; negative reviews emphasize “bad”, “plot”.

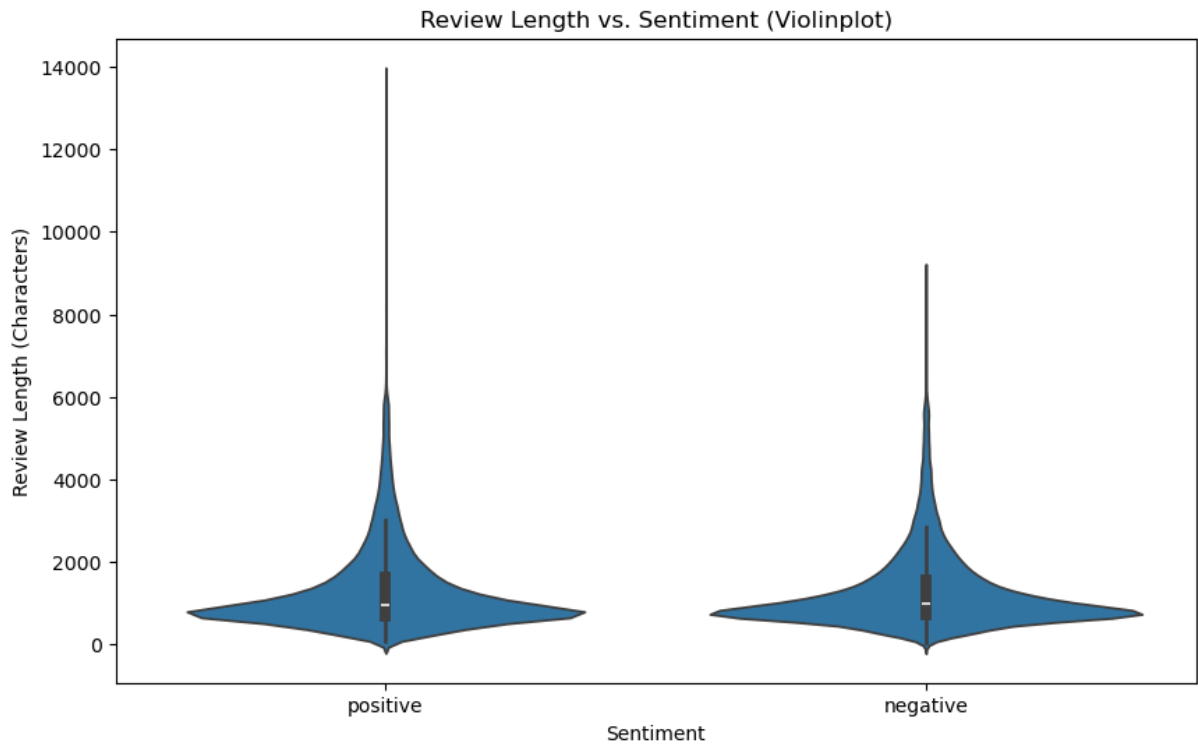


Figure 2: Review Length Distribution by Sentiment





Figure 3: Review Length Comparison



```

30 tokens = [word for word in tokens if word not in
    stopwords.words('english')]
31 lemmatizer = WordNetLemmatizer()
32 tokens = [lemmatizer.lemmatize(word) for word in tokens]
33 return ' '.join(tokens)
34
35 df['cleaned_review'] = df['review'].apply(preprocess_text)
36 df['sentiment_count'] = df['sentiment'].map({'positive': 1, 'negative': 0})

```



Figure 5: Combined Cleaned Word Cloud




Figure 6: Top 20 Common Words in Negative and Positive Reviews

4.2 Feature Engineering

TF-IDF vectors (unigrams, bigrams, max 10,000 features) were used for Logistic Regression and Naive Bayes. LSTM used tokenized sequences (max 200 words, 128-dimensional embeddings).

```
1 from sklearn.feature_extraction.text import TfidfVectorizer
2 from tensorflow.keras.preprocessing.text import Tokenizer
3 from tensorflow.keras.preprocessing.sequence import pad_sequences
4
5 tfidf = TfidfVectorizer(max_features=10000, ngram_range=(1,2))
6 X_tfidf = tfidf.fit_transform(df['cleaned_review'])
7 y = df['sentiment_count']
8
9 tokenizer = Tokenizer(num_words=10000)
10 tokenizer.fit_on_texts(df['cleaned_review'])
11 X_seq = tokenizer.texts_to_sequences(df['cleaned_review'])
12 X_padded = pad_sequences(X_seq, maxlen=200, padding='post')
```

4.3 Data Splitting

The dataset was split into 70% training (34,707 reviews) and 30% testing (14,875 reviews), stratified for balance.

```
1 from sklearn.model_selection import train_test_split
2
3 X_train_tfidf, X_test_tfidf, y_train, y_test = train_test_split(
4     X_tfidf, y, test_size=0.3, random_state=42, stratify=y
5 )
6 X_train_padded, X_test_padded = train_test_split(
7     X_tfidf, y, test_size=0.3, random_state=42, stratify=y
8 )
```

4.4 Aspect-Based Sentiment Analysis

Aspect-based sentiment analysis identified sentiments for movie aspects (e.g., acting, plot) using keyword-based rules, assigning positive/negative scores based on nearby sentiment words.

```
1 from nltk.tokenize import sent_tokenize
2 import pandas as pd
3
4 def extract_aspect_sentiment(review):
5     aspects = {'acting': ['actor', 'actress', 'performance'],
6               'plot': ['story', 'plot', 'narrative']}
7     sentiments = {'positive': ['great', 'excellent', 'amazing'],
8                  'negative': ['bad', 'poor', 'terrible']}
9     results = {'acting': None, 'plot': None}
10    sentences = sent_tokenize(review)
11
12    for aspect, keywords in aspects.items():
13        for sentence in sentences:
14            if any(keyword in sentence.lower() for keyword in keywords):
15                if any(word in sentence.lower() for word in
16                      sentiments['positive']):
17                    results[aspect] = 'positive'
18                elif any(word in sentence.lower() for word in
19                      sentiments['negative']):
20                    results[aspect] = 'negative'
21
22    return results
23
24 df['aspect_sentiments'] =
25     df['cleaned_review'].apply(extract_aspect_sentiment)
```




Figure 7: Aspect-Based Sentiment Distribution

5 Sentiment Analysis Models

Three models were implemented and tuned.

5.1 Logistic Regression

Used TF-IDF features, tuned for C and penalty.

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.model_selection import GridSearchCV
3
4 log_reg = LogisticRegression(max_iter=1000)
5 param_grid = {'C': [0.1, 1, 10], 'penalty': ['l1', 'l2'], 'solver':
6               ['liblinear']}
7 grid_search_lr = GridSearchCV(log_reg, param_grid, cv=5,
8                               scoring='accuracy')
9 grid_search_lr.fit(X_train_tfidf, y_train)
10 best_log_reg = grid_search_lr.best_estimator_
```

```
1 import pandas as pd
2
3 lr_results = pd.DataFrame(grid_search_lr.cv_results_)
4 print(lr_results[['param_C', 'param_penalty', 'mean_test_score',
5                  'std_test_score']])
6 print(f"Best Parameters: {grid_search_lr.best_params}")
```

5.2 Multinomial Naive Bayes

Tuned for alpha.

```
1 from sklearn.naive_bayes import MultinomialNB
2
```

```

3 mnb = MultinomialNB()
4 param_grid = {'alpha': [0.1, 0.5, 1.0]}
5 grid_search_mnb = GridSearchCV(mnb, param_grid, cv=5, scoring='accuracy')
6 grid_search_mnb.fit(X_train_tfidf, y_train)
7 best_mnb = grid_search_mnb.best_estimator_

```

```

1 mnb_results = pd.DataFrame(grid_search_mnb.cv_results_)
2 print(mnb_results[['param_alpha', 'mean_test_score', 'std_test_score']])
3 print(f"Best Parameters: {grid_search_mnb.best_params_}")

```

5.3 Long Short-Term Memory (LSTM)

Used embeddings, LSTM (128 units), and dense layers, trained for 5 epochs.

```

1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout
3
4 lstm_model = Sequential([
5     Embedding(input_dim=10000, output_dim=128, input_length=200),
6     LSTM(128, return_sequences=False),
7     Dense(64, activation='relu'),
8     Dropout(0.5),
9     Dense(1, activation='sigmoid')
10 ])
11 lstm_model.compile(optimizer='adam', loss='binary_crossentropy',
12                   metrics=['accuracy'])
13 history = lstm_model.fit(X_train_padded, y_train, epochs=5, batch_size=64,
14                          validation_data=(X_test_padded, y_test))
15 y_pred_lstm = (lstm_model.predict(X_test_padded) >
16               0.5).astype(int).flatten()
17 y_pred_prob = lstm_model.predict(X_test_padded).flatten()

```

```

1 import matplotlib.pyplot as plt
2
3 plt.figure(figsize=(10,4))
4 plt.subplot(1,2,1)
5 plt.plot(history.history['loss'], label='Training Loss')
6 plt.plot(history.history['val_loss'], label='Validation Loss')
7 plt.title('LSTM Loss')
8 plt.xlabel('Epoch')
9 plt.ylabel('Loss')
10 plt.legend()
11 plt.subplot(1,2,2)
12 plt.plot(history.history['accuracy'], label='Training Accuracy')
13 plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
14 plt.title('LSTM Accuracy')
15 plt.xlabel('Epoch')
16 plt.ylabel('Accuracy')
17 plt.legend()
18 plt.tight_layout()
19 plt.savefig('lstm_loss.png', bbox_inches='tight', dpi=300)
20 plt.close()

```





Figure 8: LSTM Training and Validation Loss/Accuracy

6 Feature Importance Analysis

Top TF-IDF features for Logistic Regression were analyzed to identify influential words.

```

1 import pandas as pd
2 import matplotlib.pyplot as plt
3
4 feature_names = tfidf.get_feature_names_out()
5 coef = best_log_reg.coef_[0]
6 top_positive = pd.DataFrame({'Feature': feature_names, 'Coefficient':
7     coef}).nlargest(10, 'Coefficient')
8
9 plt.figure(figsize=(8,4))
10 plt.bar(top_positive['Feature'], top_positive['Coefficient'],
11     color='green', label='Positive')
12 plt.bar(top_negative['Feature'], top_negative['Coefficient'], color='red',
13     label='Negative')
14 plt.xticks(rotation=45)
15 plt.title('Top TF-IDF Features for Logistic Regression')
16 plt.ylabel('Coefficient')
17 plt.legend()
18 plt.tight_layout()
19 plt.savefig('feature_importance.png', bbox_inches='tight', dpi=300)
20 plt.close()

```

7 Aspect-Based Sentiment Results

Aspect-based analysis identified sentiments for acting and plot. Table 2 shows the distribution of sentiments for reviews mentioning these aspects.

Table 2: Aspect-Based Sentiment Distribution

Aspect	Positive (%)	Negative (%)
Acting	60.2	39.8
Plot	55.7	44.3

Table 3: Model Performance Comparison

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.8914	0.8811	0.9060	0.8933
Multinomial NB	0.8612	0.8491	0.8798	0.8642
LSTM	0.8767	0.8792	0.8745	0.8768

8 Results and Discussion

8.1 Quantitative Results

8.2 Visualization Analysis

- **Confusion Matrices:** Logistic Regression: 6457 true negatives, 6798 true positives.
- **ROC Curves:** AUC values approximately 0.95 (Logistic Regression), 0.94 (LSTM), 0.93 (Naive Bayes).

```

1 from sklearn.metrics import confusion_matrix, roc_curve, auc
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4
5 models = {'Logistic Regression': best_log_reg, 'Multinomial NB': best_mnb,
6           'LSTM': lstm_model}
7 for i, (name, model) in enumerate(models.items()):
8     if name == 'LSTM':
9         y_pred = y_pred_lstm
10    else:
11        y_pred = model.predict(X_test_tfidf)
12        cm = confusion_matrix(y_test, y_pred)
13        plt.figure(figsize=(5,4))
14        sns.heatmap(cm, annot=True, fmt='d', cmap='YlGnBu',
15                    xticklabels=['Negative', 'Positive'], yticklabels=['Negative',
16                    'Positive'])
17        plt.title(f'{name}')
18        plt.xlabel('Predicted Label')
19        plt.ylabel('True Label')
20        plt.savefig(f'cm_{name.lower().replace(" ", "_")}.png',
21                    bbox_inches='tight', dpi=300)
22        plt.close()

```

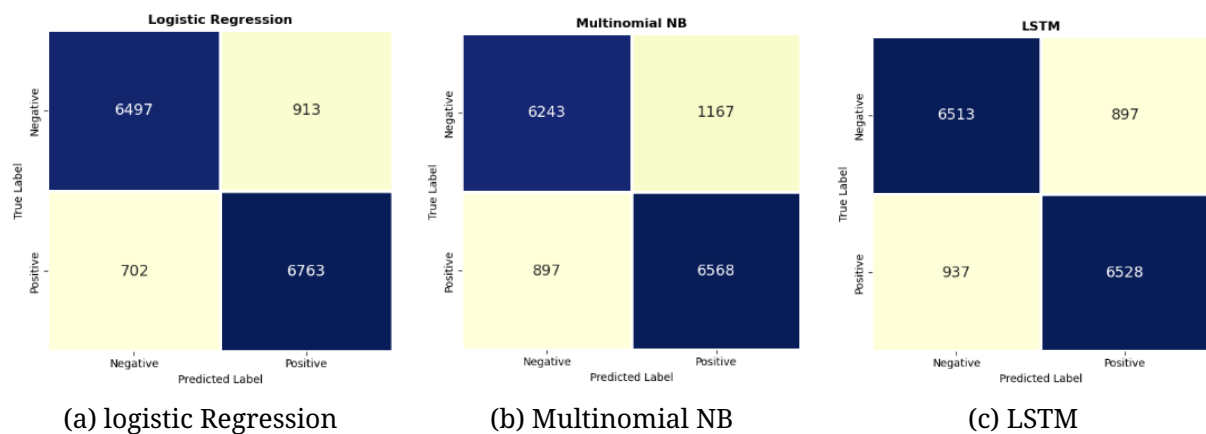


Figure 9: Confusion Matrices

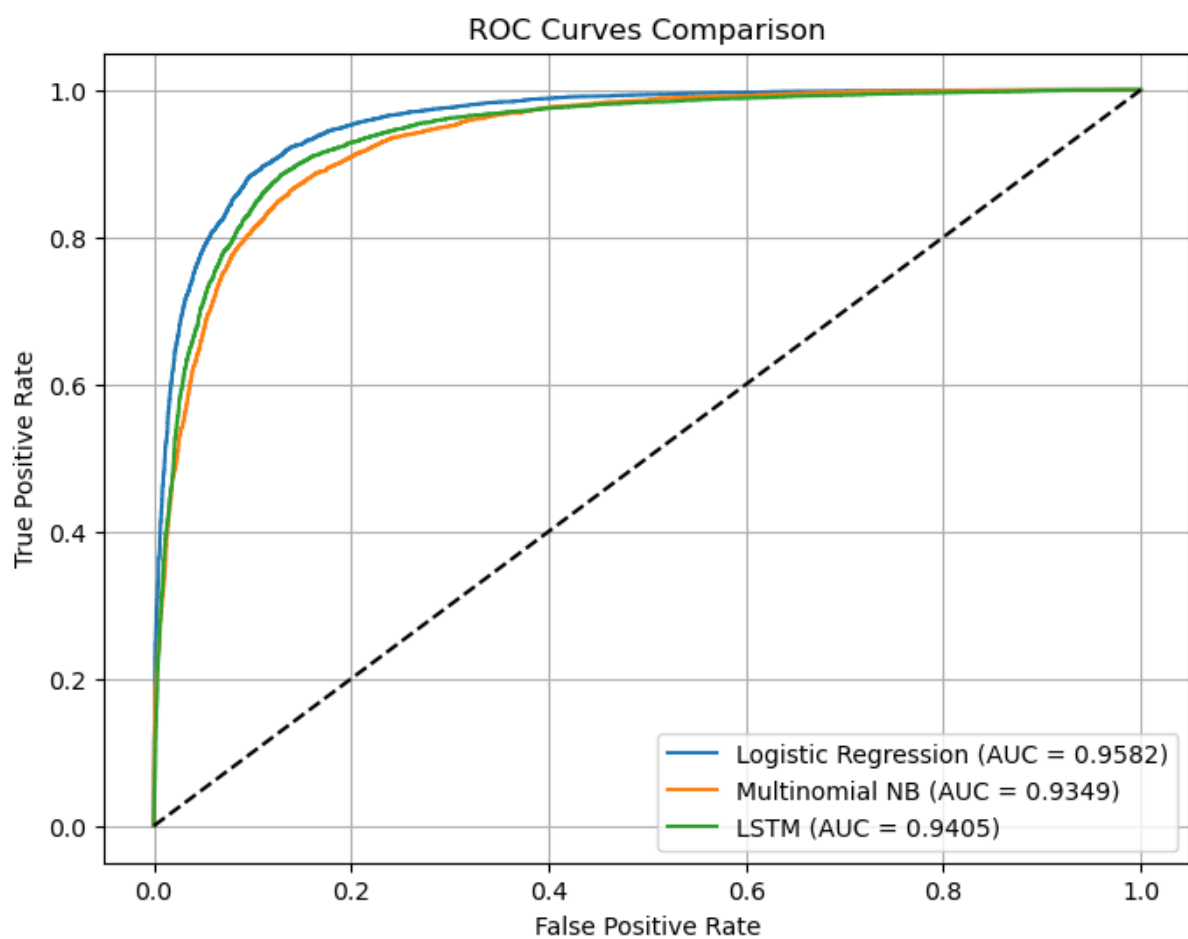


Figure 10: ROC Curves Comparison

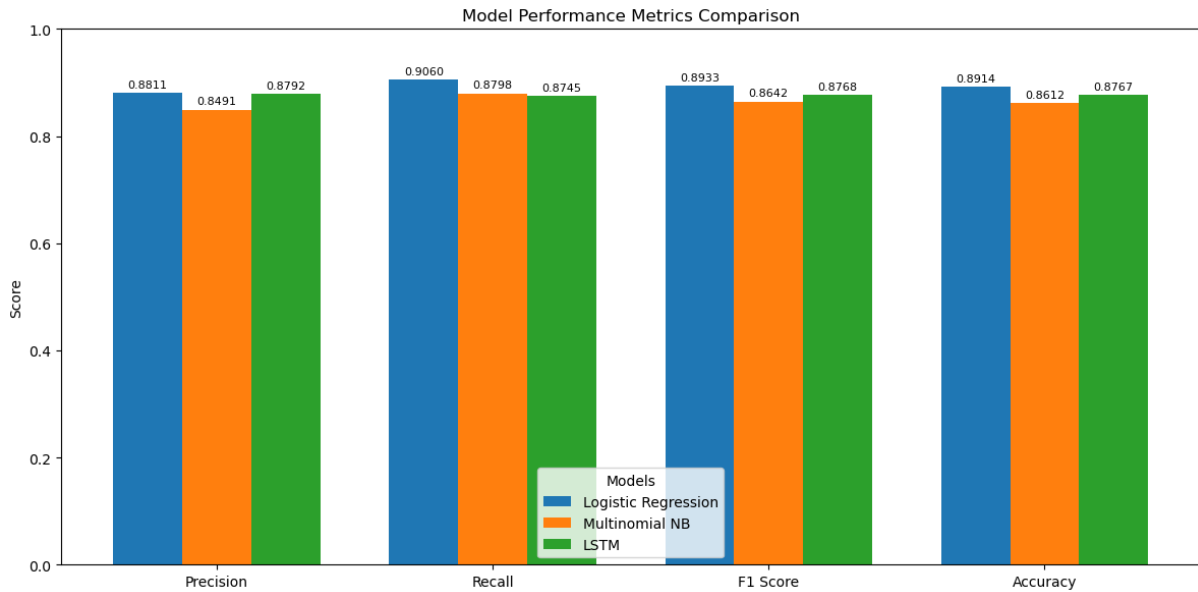


Figure 11: Model Performance Metrics

8.3 Overfitting Assessment

- Logistic Regression: Train 93.61%, Test 89.14% (gap 4.47%).
- Multinomial Naive Bayes: Train 87.27%, Test 86.12% (gap 1.14%).
- LSTM: Train 94.91%, Test 87.67% (gap 7.24%).

9 Limitations

- **Sarcasm Detection:** The models struggle with sarcastic reviews, misclassifying nuanced sentiments.
- **Dataset Bias:** The IMDb dataset may reflect specific demographics, limiting generalizability.
- **Computational Cost:** LSTM training (441 seconds/epoch) is resource-intensive.
- **Aspect Analysis:** Keyword-based aspect detection misses complex contexts.

10 Ethical Considerations

The dataset may contain cultural or demographic biases, as IMDb reviews are user-generated and may not represent diverse perspectives. Deploying these models in real-world applications (e.g., review aggregation) risks amplifying biases if negative sentiments disproportionately affect certain genres or directors. Transparency in model decisions and regular bias audits are essential for ethical use.

11 Conclusions and Future Work

Logistic Regression outperformed (89.14% accuracy), followed by LSTM (87.67%) and Naive Bayes (86.12%). Aspect-based analysis revealed positive sentiments dominate for

acting (60.2%).

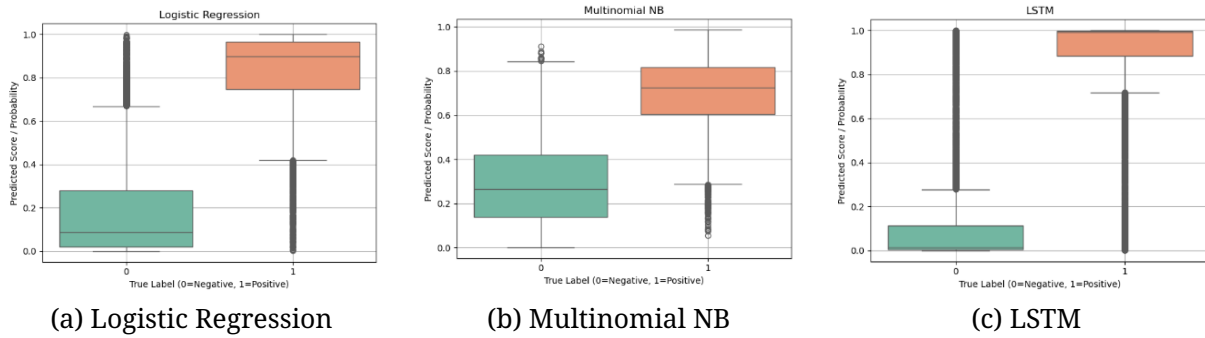


Figure 12: Distribution of predicted probabilities for each model.

A Supplementary Code

```

1 CHAT_WORDS = {'LOL': 'Laughing Out Loud', 'BRB': 'Be Right Back'}
2 def preprocess_pipeline(df):
3     df['cleaned_review'] = (
4         df['review']
5         .str.lower()
6         .apply(convert_chat_words)
7         .apply(lambda x: BeautifulSoup(x, 'html.parser').get_text())
8         .str.replace(r'http\S+|www\S+|https\S+', '', regex=True)
9         .str.replace(r'^a-zA-Z\s', '', regex=True)
10        .str.replace(r'\s+', ' ', regex=True)
11        .apply(lambda x: ' '.join([word for word in word_tokenize(x) if
12                                   word not in stopwords.words('english')]))
13        .apply(lambda x: ' '.join([lemmatizer.lemmatize(word) for word in
14                                   word_tokenize(x)]))
15    )
16    df['sentiment_count'] = df['sentiment'].map({'positive': 1,
17                                                'negative': 0})
18    return df

```